

## IMPORT DATASET

```
In [1]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
import warnings  
warnings.filterwarnings('ignore')
```

## LOADING THE CSV FILE (DATASET)

```
In [2]: df= pd.read_csv('Heart_Attack_prediction_Dataset.csv')  
df.head()
```

Out[2]:

	Patient ID	Age	Sex	Cholesterol	Blood Pressure	Heart Rate	Diabetes	Family History	Smoking	Obesity	...	Sedentary Hours Per Day	Income	BI
0	BMW7812	67	Male	208	158/88	72	0	0	1	0	...	6.615001	261404	31.2512
1	CZE1114	21	Male	389	165/93	98	1	1	1	1	...	4.963459	285768	27.1949
2	BNI9906	21	Female	324	174/99	72	1	0	0	0	...	9.463426	235282	28.1765
3	JLN3497	84	Male	383	163/100	73	1	1	1	0	...	7.648981	125640	36.4647
4	GFO8847	66	Male	318	91/88	93	1	1	1	1	...	1.514821	160555	21.8091

5 rows × 26 columns

In [3]: `df.columns`

```
Out[3]: Index(['Patient ID', 'Age', 'Sex', 'Cholesterol', 'Blood Pressure',
       'Heart Rate', 'Diabetes', 'Family History', 'Smoking', 'Obesity',
       'Alcohol Consumption', 'Exercise Hours Per Week', 'Diet',
       'Previous Heart Problems', 'Medication Use', 'Stress Level',
       'Sedentary Hours Per Day', 'Income', 'BMI', 'Triglycerides',
       'Physical Activity Days Per Week', 'Sleep Hours Per Day', 'Country',
       'Continent', 'Hemisphere', 'Heart Attack Risk'],
      dtype='object')
```

In [4]: `df.drop_duplicates(inplace=True)`  
`df.shape`

Out[4]: (8763, 26)

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8763 entries, 0 to 8762
Data columns (total 26 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Patient ID      8763 non-null   object  
 1   Age              8763 non-null   int64  
 2   Sex              8763 non-null   object  
 3   Cholesterol     8763 non-null   int64  
 4   Blood Pressure   8763 non-null   object  
 5   Heart Rate       8763 non-null   int64  
 6   Diabetes         8763 non-null   int64  
 7   Family History   8763 non-null   int64  
 8   Smoking          8763 non-null   int64  
 9   Obesity          8763 non-null   int64  
 10  Alcohol Consumption 8763 non-null   int64  
 11  Exercise Hours Per Week 8763 non-null   float64 
 12  Diet              8763 non-null   object  
 13  Previous Heart Problems 8763 non-null   int64  
 14  Medication Use   8763 non-null   int64  
 15  Stress Level     8763 non-null   int64  
 16  Sedentary Hours Per Day 8763 non-null   float64 
 17  Income            8763 non-null   int64  
 18  BMI               8763 non-null   float64 
 19  Triglycerides    8763 non-null   int64  
 20  Physical Activity Days Per Week 8763 non-null   int64  
 21  Sleep Hours Per Day 8763 non-null   int64  
 22  Country           8763 non-null   object  
 23  Continent         8763 non-null   object  
 24  Hemisphere        8763 non-null   object  
 25  Heart Attack Risk 8763 non-null   int64  
dtypes: float64(3), int64(16), object(7)
memory usage: 1.7+ MB
```

```
In [6]: df.describe().T
```

## Heart Disease Attack

Out[6]:		count	mean	std	min	25%	50%	75%	max
	<b>Age</b>	8763.0	53.707977	21.249509	18.000000	35.000000	54.000000	72.000000	90.000000
	<b>Cholesterol</b>	8763.0	259.877211	80.863276	120.000000	192.000000	259.000000	330.000000	400.000000
	<b>Heart Rate</b>	8763.0	75.021682	20.550948	40.000000	57.000000	75.000000	93.000000	110.000000
	<b>Diabetes</b>	8763.0	0.652288	0.476271	0.000000	0.000000	1.000000	1.000000	1.000000
	<b>Family History</b>	8763.0	0.492982	0.499979	0.000000	0.000000	0.000000	1.000000	1.000000
	<b>Smoking</b>	8763.0	0.896839	0.304186	0.000000	1.000000	1.000000	1.000000	1.000000
	<b>Obesity</b>	8763.0	0.501426	0.500026	0.000000	0.000000	1.000000	1.000000	1.000000
	<b>Alcohol Consumption</b>	8763.0	0.598083	0.490313	0.000000	0.000000	1.000000	1.000000	1.000000
	<b>Exercise Hours Per Week</b>	8763.0	10.014284	5.783745	0.002442	4.981579	10.069559	15.050018	19.998709
	<b>Previous Heart Problems</b>	8763.0	0.495835	0.500011	0.000000	0.000000	0.000000	1.000000	1.000000
	<b>Medication Use</b>	8763.0	0.498345	0.500026	0.000000	0.000000	0.000000	1.000000	1.000000
	<b>Stress Level</b>	8763.0	5.469702	2.859622	1.000000	3.000000	5.000000	8.000000	10.000000
	<b>Sedentary Hours Per Day</b>	8763.0	5.993690	3.466359	0.001263	2.998794	5.933622	9.019124	11.999313
	<b>Income</b>	8763.0	158263.181901	80575.190806	20062.000000	88310.000000	157866.000000	227749.000000	299954.000000
	<b>BMI</b>	8763.0	28.891446	6.319181	18.002337	23.422985	28.768999	34.324594	39.997211
	<b>Triglycerides</b>	8763.0	417.677051	223.748137	30.000000	225.500000	417.000000	612.000000	800.000000
	<b>Physical Activity Days</b>	8763.0	3.489672	2.282687	0.000000	2.000000	3.000000	5.000000	7.000000

	count	mean	std	min	25%	50%	75%	max
<b>Per Week</b>								
Sleep Hours Per Day	8763.0	7.023508	1.988473	4.000000	5.000000	7.000000	9.000000	10.000000
Heart Attack Risk	8763.0	0.358211	0.479502	0.000000	0.000000	0.000000	1.000000	1.000000

## DATA CLEANING

```
In [7]: columns_to_drop = ['Hemisphere', 'Patient ID', 'Continent']
df.drop(columns_to_drop, axis=1, inplace=True)
df.head()
```

Out[7]:

	Age	Sex	Cholesterol	Blood Pressure	Heart Rate	Diabetes	Family History	Smoking	Obesity	Alcohol Consumption	...	Medication Use	Stress Level	Sedentary Hours
0	67	Male	208	158/88	72	0	0	1	0	0	...	0	9	6.6
1	21	Male	389	165/93	98	1	1	1	1	1	...	0	1	4.9
2	21	Female	324	174/99	72	1	0	0	0	0	...	1	9	9.4
3	84	Male	383	163/100	73	1	1	1	0	1	...	0	9	7.6
4	66	Male	318	91/88	93	1	1	1	1	0	...	0	6	1.5

5 rows × 23 columns



```
In [8]: # Split 'Blood Pressure' into systolic/diastolic
if "Blood Pressure" in df.columns:
    bp_split = df["Blood Pressure"].str.split("/", expand=True)
    df["Systolic_BP"] = pd.to_numeric(bp_split[0], errors="coerce")
    df["Diastolic_BP"] = pd.to_numeric(bp_split[1], errors="coerce")
```

```
df.drop("Blood Pressure", axis=1, inplace=True)

# Confirm final structure
print("Cleaning, rounding & normalization complete!")
print(f"Final shape: {df.shape}")
print("Data preview:")
print(df.head())
```

Cleaning, rounding & normalization complete!

Final shape: (8763, 24)

Data preview:

	Age	Sex	Cholesterol	Heart Rate	Diabetes	Family History	Smoking	\
0	67	Male	208	72	0	0	1	
1	21	Male	389	98	1	1	1	
2	21	Female	324	72	1	0	0	
3	84	Male	383	73	1	1	1	
4	66	Male	318	93	1	1	1	

	Obesity	Alcohol Consumption	Exercise Hours Per Week	...	\
0	0	0	4.168189	...	
1	1	1	1.813242	...	
2	0	0	2.078353	...	
3	0	1	9.828130	...	
4	1	0	5.804299	...	

	Sedentary Hours Per Day	Income	BMI	Triglycerides	\
0	6.615001	261404	31.251233	286	
1	4.963459	285768	27.194973	235	
2	9.463426	235282	28.176571	587	
3	7.648981	125640	36.464704	378	
4	1.514821	160555	21.809144	231	

	Physical Activity Days Per Week	Sleep Hours Per Day	Country	\
0	0	6	Argentina	
1	1	7	Canada	
2	4	4	France	
3	3	4	Canada	
4	1	5	Thailand	

	Heart Attack Risk	Systolic_BP	Diastolic_BP	
0	0	158	88	
1	0	165	93	
2	0	174	99	
3	0	163	100	
4	0	91	88	

[5 rows x 24 columns]

In [9]:

```
#Check for missing values
missing_summary =df.isnull().sum()
```

```
total_missing = missing_summary.sum()

print("\nMissing values per column:")
print(missing_summary)

print("\nData types summary after cleaning:")
print(df.info())
```

Missing values per column:

```
Age          0
Sex          0
Cholesterol 0
Heart Rate   0
Diabetes     0
Family History 0
Smoking      0
Obesity      0
Alcohol Consumption 0
Exercise Hours Per Week 0
Diet          0
Previous Heart Problems 0
Medication Use 0
Stress Level  0
Sedentary Hours Per Day 0
Income         0
BMI           0
Triglycerides 0
Physical Activity Days Per Week 0
Sleep Hours Per Day 0
Country        0
Heart Attack Risk 0
Systolic_BP    0
Diastolic_BP   0
dtype: int64
```

Data types summary after cleaning:

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 8763 entries, 0 to 8762
```

Data columns (total 24 columns):

#	Column	Non-Null Count	Dtype
0	Age	8763 non-null	int64
1	Sex	8763 non-null	object
2	Cholesterol	8763 non-null	int64
3	Heart Rate	8763 non-null	int64
4	Diabetes	8763 non-null	int64
5	Family History	8763 non-null	int64
6	Smoking	8763 non-null	int64
7	Obesity	8763 non-null	int64
8	Alcohol Consumption	8763 non-null	int64

```
9   Exercise Hours Per Week      8763 non-null  float64
10  Diet                         8763 non-null  object
11  Previous Heart Problems     8763 non-null  int64
12  Medication Use              8763 non-null  int64
13  Stress Level                8763 non-null  int64
14  Sedentary Hours Per Day    8763 non-null  float64
15  Income                       8763 non-null  int64
16  BMI                          8763 non-null  float64
17  Triglycerides               8763 non-null  int64
18  Physical Activity Days Per Week 8763 non-null  int64
19  Sleep Hours Per Day         8763 non-null  int64
20  Country                      8763 non-null  object
21  Heart Attack Risk            8763 non-null  int64
22  Systolic_BP                 8763 non-null  int64
23  Diastolic_BP                8763 non-null  int64
dtypes: float64(3), int64(18), object(3)
memory usage: 1.6+ MB
None
```

## FEATURE ENGINEERING

```
In [10]: # Feature Engineering
df["BMI_Stress"] = df["BMI"] * df["Stress Level"]
df["Activity_Ratio"] = df["Exercise Hours Per Week"] / (df["Sedentary Hours Per Day"] + 1)
df["BP_Product"] = df["Systolic_BP"] * df["Diastolic_BP"]
df["Sleep_Stress_Interaction"] = df["Sleep Hours Per Day"] * df["Stress Level"]
df["Substance_Use"] = df["Smoking"] * df["Alcohol Consumption"]
```

```
In [11]: print("\nData types summary after cleaning:")
print(df.info())
```

Data types summary after cleaning:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 8763 entries, 0 to 8762

Data columns (total 29 columns):

#	Column	Non-Null Count	Dtype
0	Age	8763 non-null	int64
1	Sex	8763 non-null	object
2	Cholesterol	8763 non-null	int64
3	Heart Rate	8763 non-null	int64
4	Diabetes	8763 non-null	int64
5	Family History	8763 non-null	int64
6	Smoking	8763 non-null	int64
7	Obesity	8763 non-null	int64
8	Alcohol Consumption	8763 non-null	int64
9	Exercise Hours Per Week	8763 non-null	float64
10	Diet	8763 non-null	object
11	Previous Heart Problems	8763 non-null	int64
12	Medication Use	8763 non-null	int64
13	Stress Level	8763 non-null	int64
14	Sedentary Hours Per Day	8763 non-null	float64
15	Income	8763 non-null	int64
16	BMI	8763 non-null	float64
17	Triglycerides	8763 non-null	int64
18	Physical Activity Days Per Week	8763 non-null	int64
19	Sleep Hours Per Day	8763 non-null	int64
20	Country	8763 non-null	object
21	Heart Attack Risk	8763 non-null	int64
22	Systolic_BP	8763 non-null	int64
23	Diastolic_BP	8763 non-null	int64
24	BMI_Stress	8763 non-null	float64
25	Activity_Ratio	8763 non-null	float64
26	BP_Product	8763 non-null	int64
27	Sleep_Stress_Interaction	8763 non-null	int64
28	Substance_Use	8763 non-null	int64

dtypes: float64(5), int64(21), object(3)

memory usage: 1.9+ MB

None

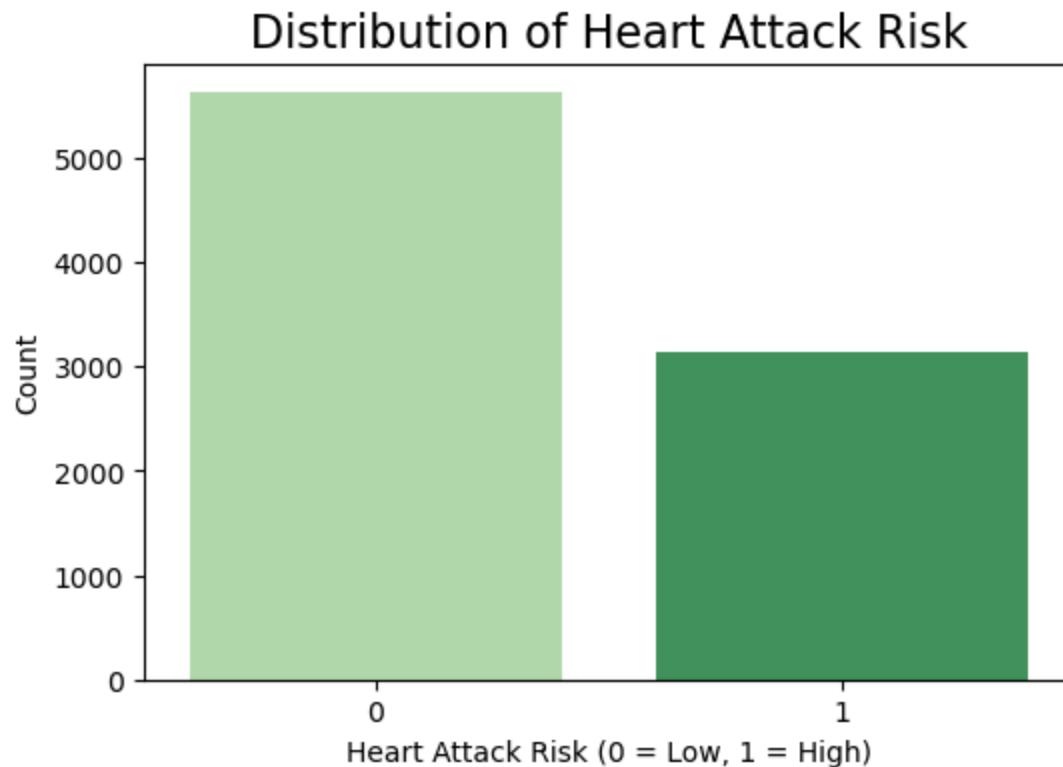
In [12]: df.unique()

```
Out[12]: Age           73  
Sex            2  
Cholesterol    281  
Heart Rate      71  
Diabetes        2  
Family History   2  
Smoking         2  
Obesity         2  
Alcohol Consumption  2  
Exercise Hours Per Week 8763  
Diet            3  
Previous Heart Problems 2  
Medication Use   2  
Stress Level     10  
Sedentary Hours Per Day 8763  
Income          8615  
BMI             8763  
Triglycerides    771  
Physical Activity Days Per Week 8  
Sleep Hours Per Day    7  
Country          20  
Heart Attack Risk   2  
Systolic_BP       91  
Diastolic_BP      51  
BMI_Stress        8763  
Activity_Ratio    8763  
BP_Product        2835  
Sleep_Stress_Interaction 39  
Substance_Use     2  
dtype: int64
```

## EXPLORATORY DATA ANALYSIS AND VISUALIZINGS

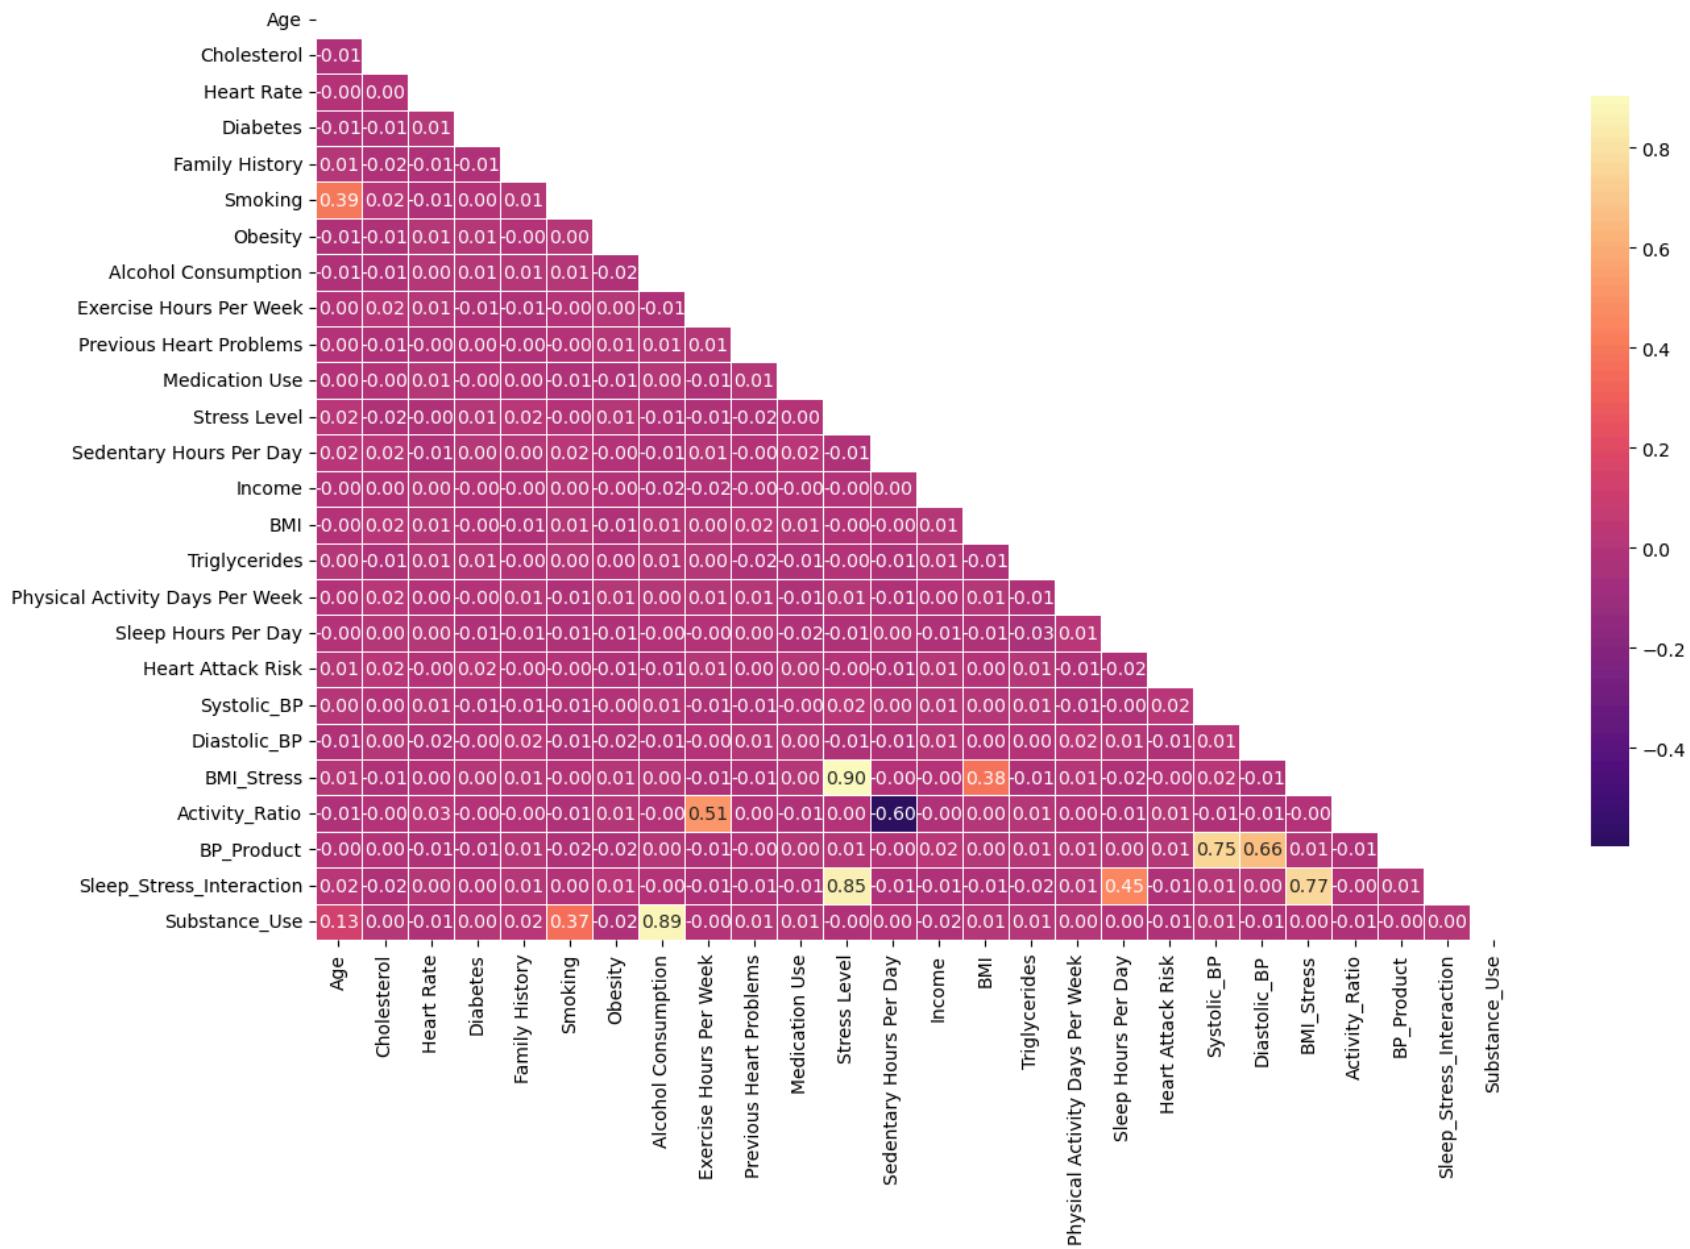
```
In [13]: # --- Distribution of Target Variable --  
plt.figure(figsize=(6,4))  
sns.countplot(x="Heart Attack Risk", hue="Heart Attack Risk", data=df, palette="Greens",  
legend=False  
)  
plt.title("Distribution of Heart Attack Risk", fontsize = 16)  
plt.xlabel("Heart Attack Risk (0 = Low, 1 = High)")
```

```
plt.ylabel("Count")
plt.show()
```



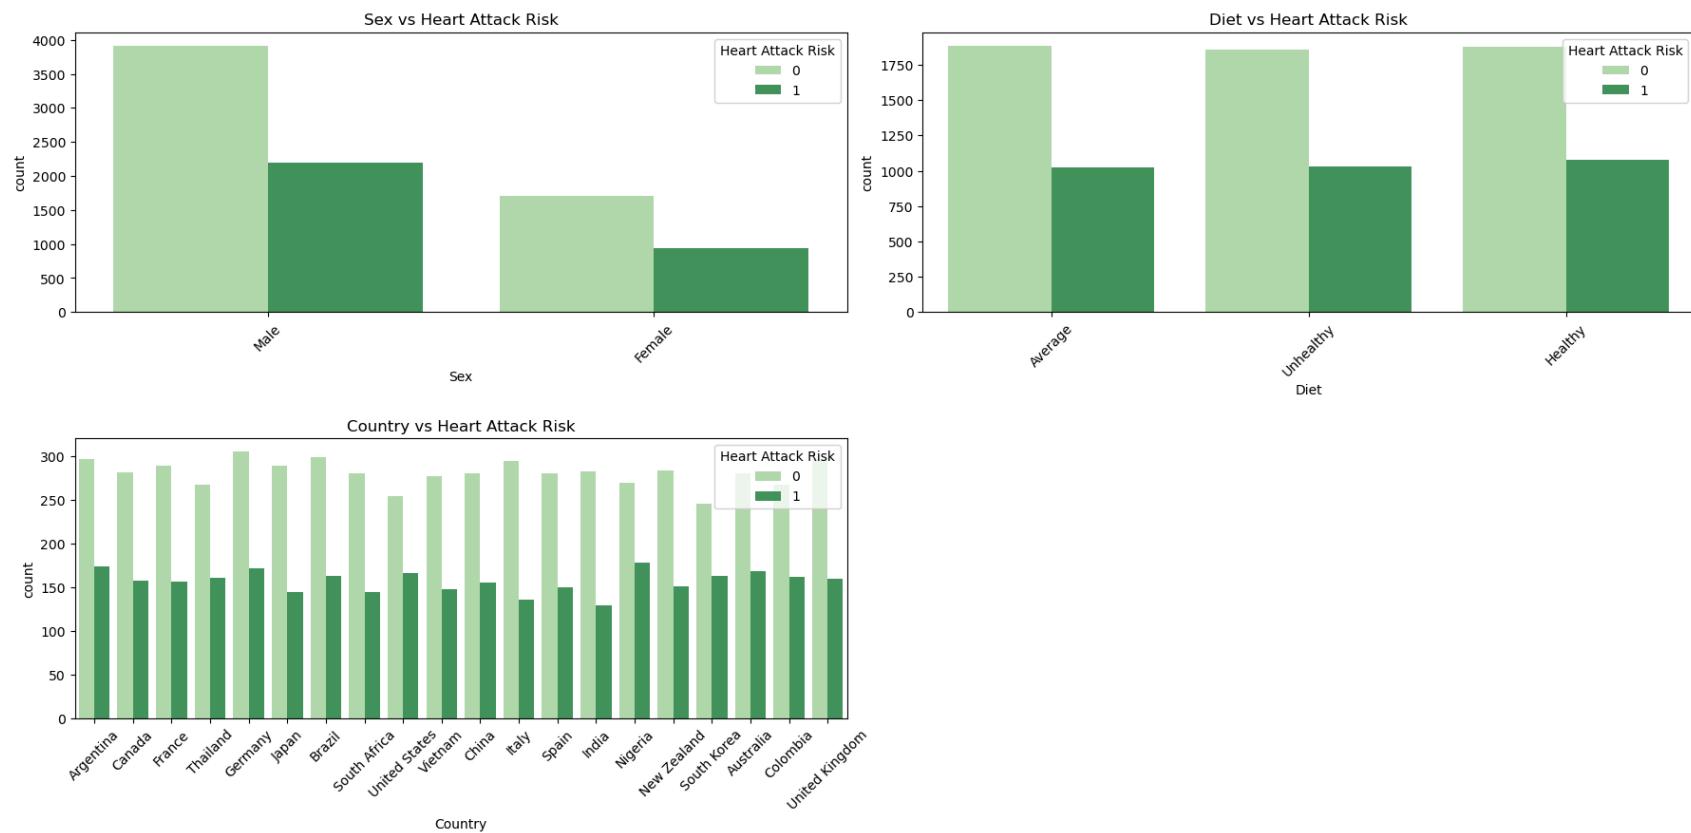
```
In [14]: # --- Correlation Heatmap --
plt.figure(figsize=(14,10))
# Select only numeric columns for correlation
corr = df.select_dtypes(include=['number']).corr()
# Focus on stronger correlations only
mask = np.triu(np.ones_like(corr, dtype=bool)) # mask upper triangle
sns.heatmap(corr, mask=mask, cmap="magma", center=0, annot=True, fmt=".2f", linewidths=0.5,
cbar_kws={"shrink": 0.8})
plt.title("Correlation HeatMap for the Numeric Features of Heart Disease Attack", fontsize=16)
plt.tight_layout()
plt.show()
```

## Correlation HeatMap for the Numeric Features of Heart Disease Attack



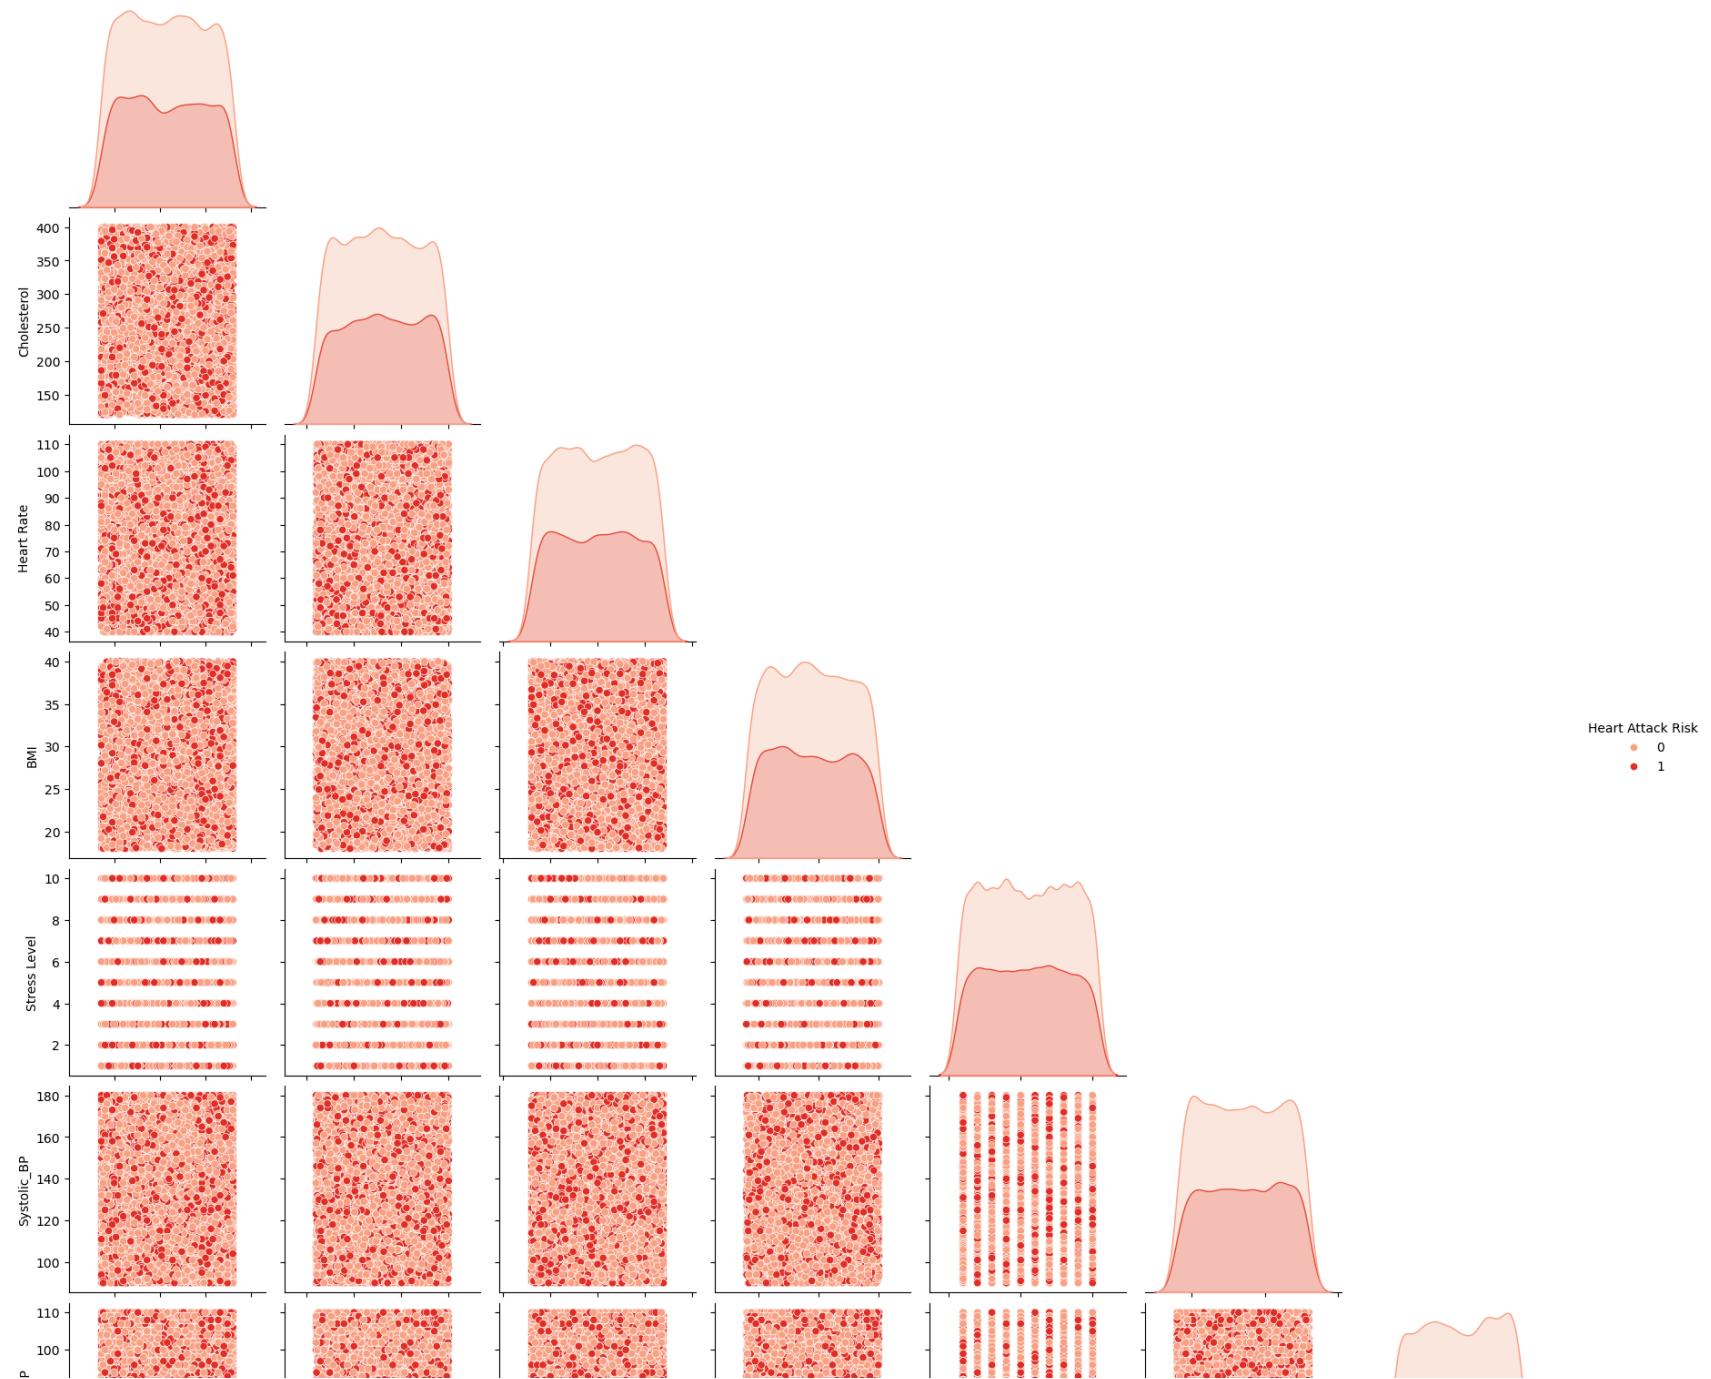
```
In [15]: # Target Distribution by Category
cat_cols = ['Sex', 'Diet', 'Country']
```

```
plt.figure(figsize=(18,12))
for i, col in enumerate(cat_cols, 1):
    plt.subplot(3, 2, i)
    sns.countplot(data=df, x=col, hue="Heart Attack Risk", palette="Greens")
    plt.title(f"{col} vs Heart Attack Risk")
    plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

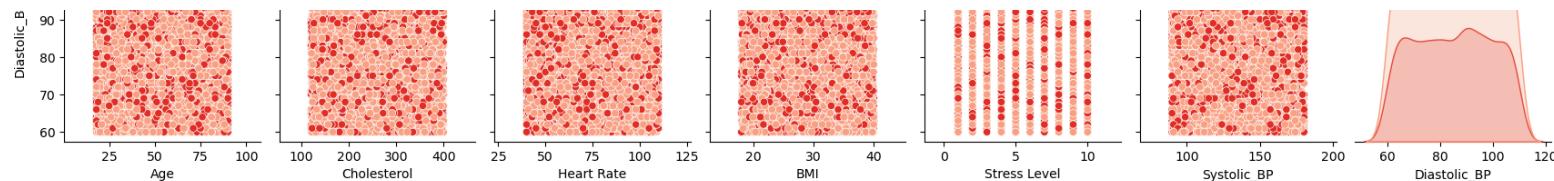


```
In [16]: # Pairwise Relationships (Hidden Correlations)
sns.pairplot( df[['Age', 'Cholesterol', 'Heart Rate', 'BMI', 'Stress Level', 'Systolic_BP', 'Diastolic_BP', 'Heart Attack Risk']], hue="Heart Attack Risk", palette="Reds", diag_kind="kde", corner=True)
plt.suptitle("Pairwise Feature Relationships", y=1.02)
plt.show()
```

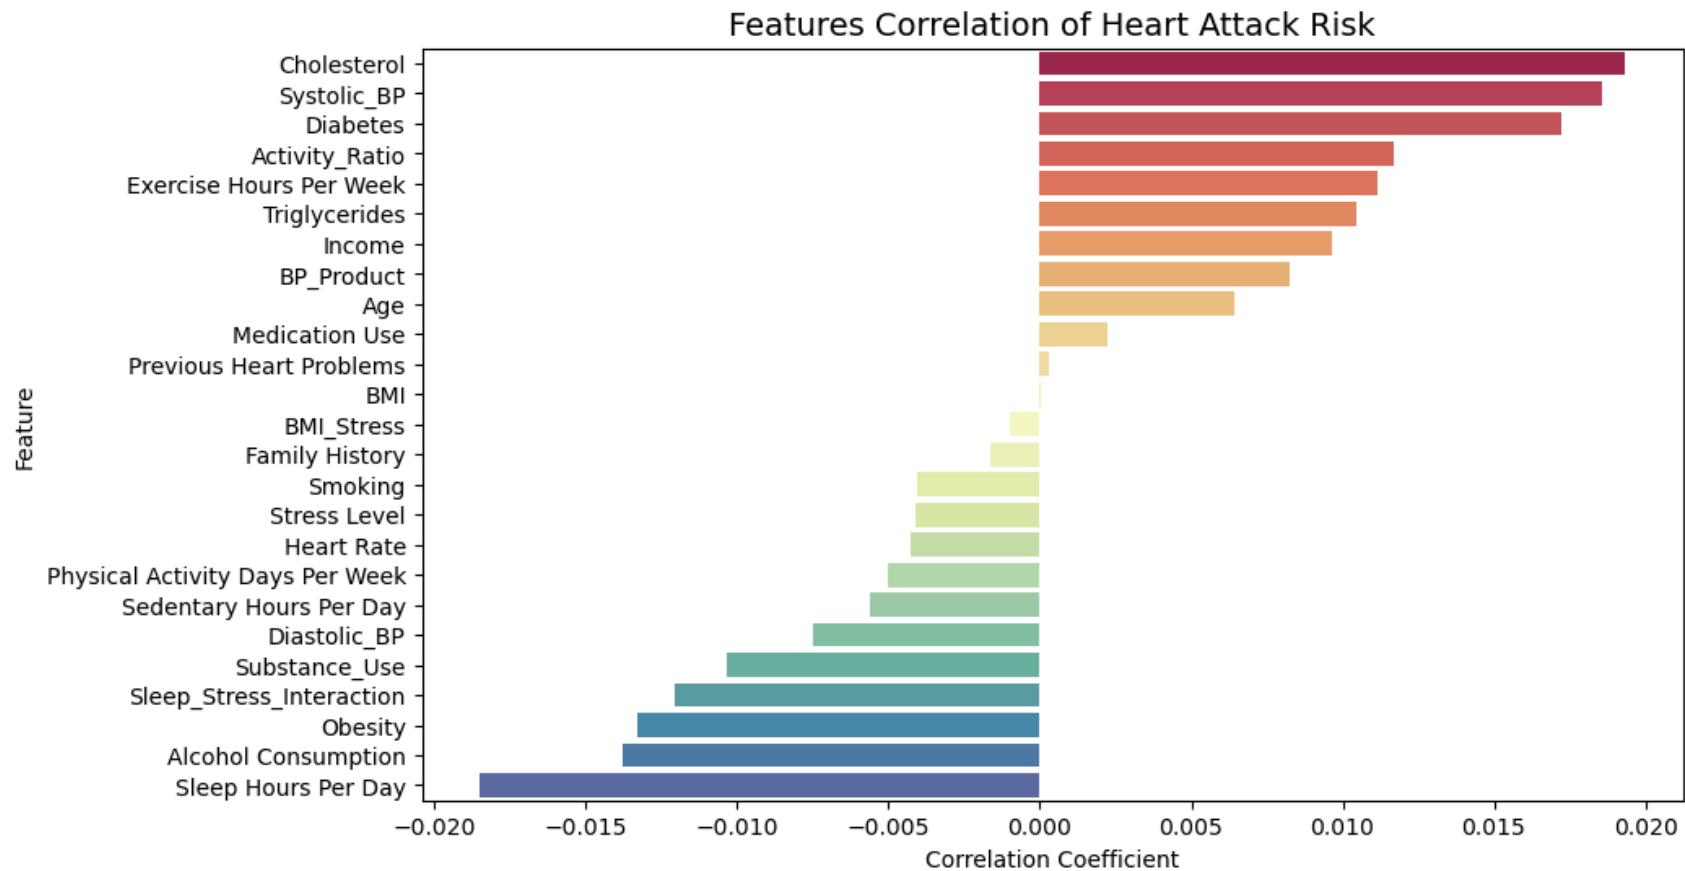
Pairwise Feature Relationships



## Heart Disease Attack



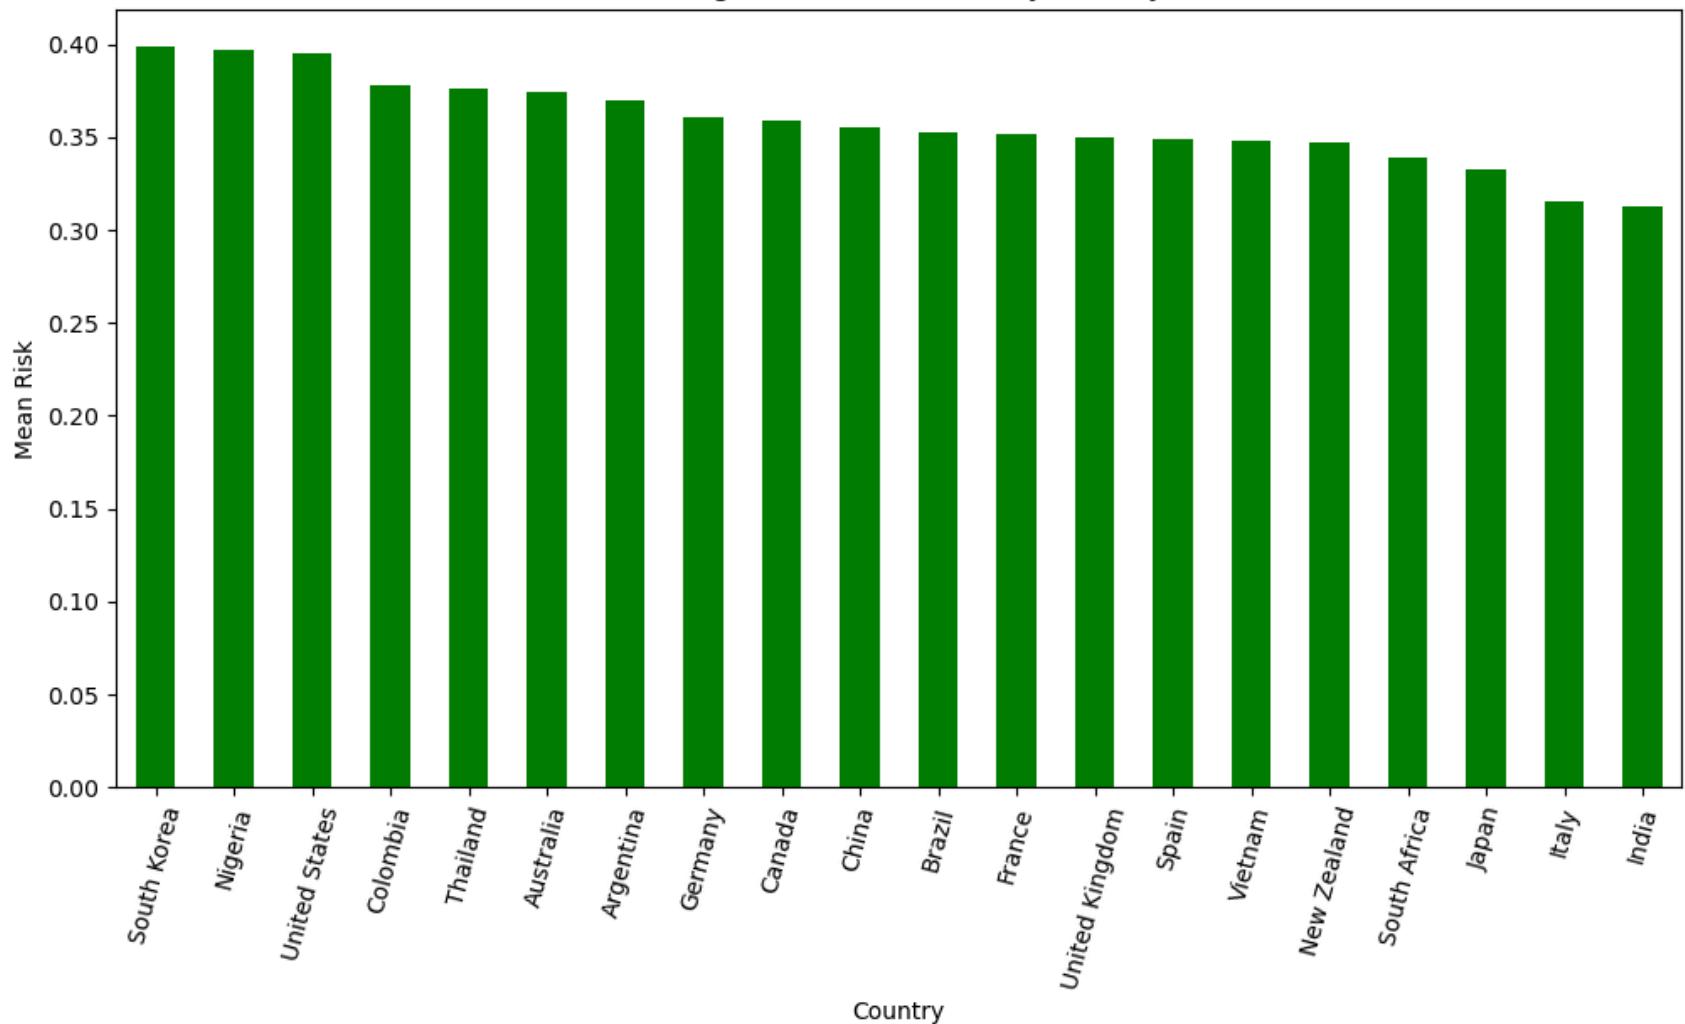
```
In [17]: # Sorted Correlation with Target
numeric_df = df.select_dtypes(include=['number'])
target_corr = (numeric_df.corr()['Heart Attack Risk']
    .drop('Heart Attack Risk')
    .sort_values(ascending=False)
)
# --- Visualize correlation strength --
plt.figure(figsize=(10, 6))
sns.barplot(x=target_corr, y=target_corr.index, palette="Spectral")
plt.title("Features Correlation of Heart Attack Risk", fontsize=14)
plt.xlabel("Correlation Coefficient")
plt.ylabel("Feature")
plt.show()
```



In [18]:

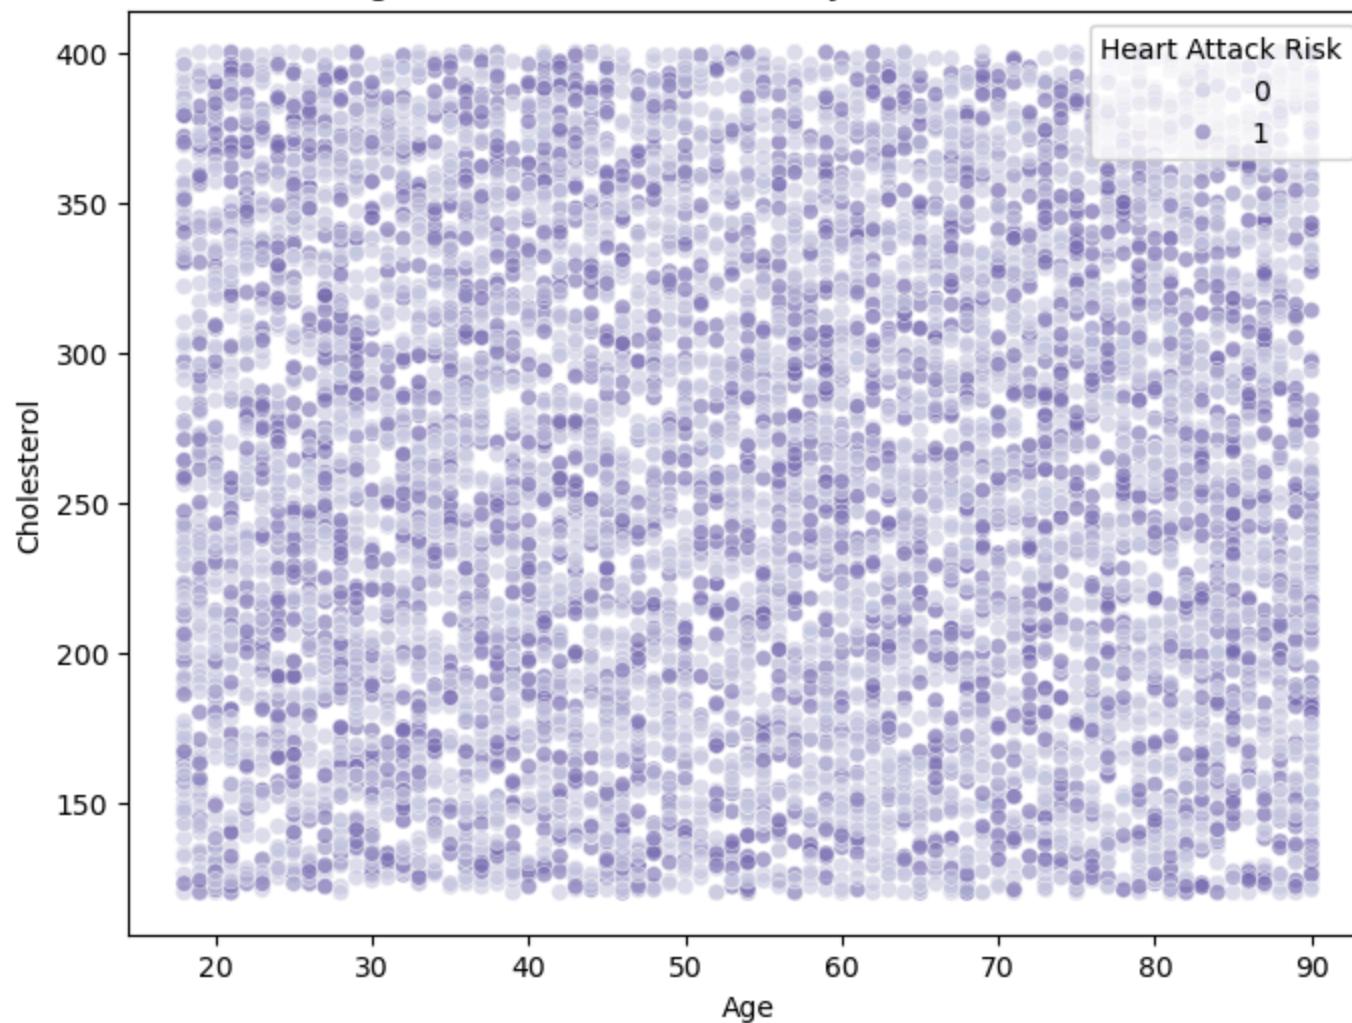
```
# Geographic Patterns(country)
risk_by_country = df.groupby('Country')['Heart Attack Risk'].mean().sort_values(ascending=False)
plt.figure(figsize=(12,6))
risk_by_country.plot(kind='bar', color='green')
plt.title("Average Heart Attack Risk by Country")
plt.ylabel("Mean Risk")
plt.xticks(rotation=75)
plt.show()
```

## Average Heart Attack Risk by Country



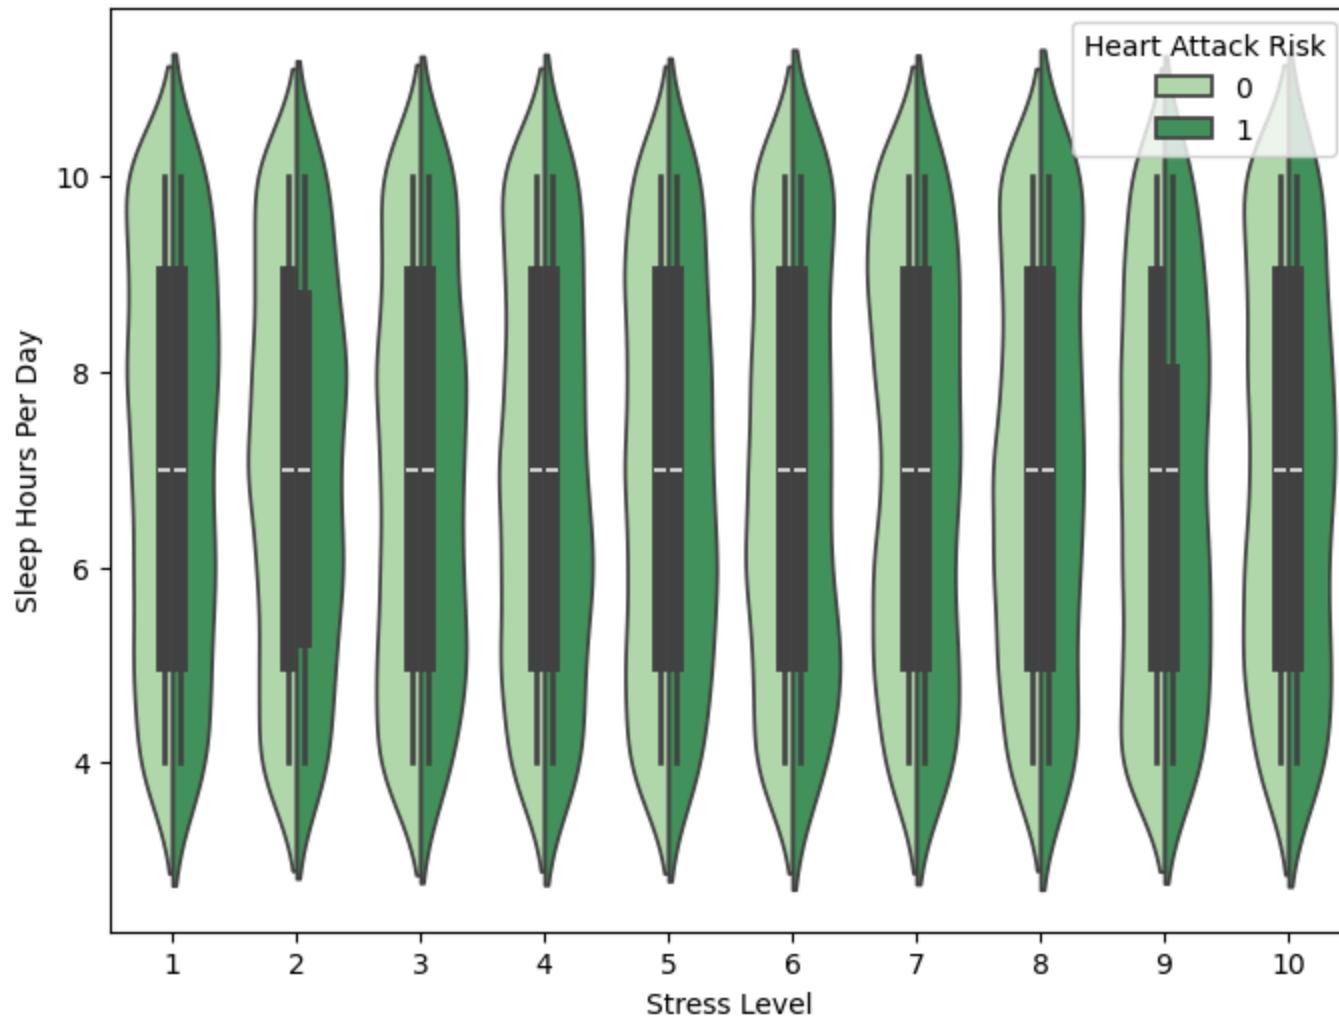
```
In [19]: # Age vs Cholesterol Colored by Heart Attack Risk
plt.figure(figsize=(8,6))
sns.scatterplot(data= df, x='Age', y='Cholesterol', hue='Heart Attack Risk', palette='Purples', alpha=0.6
)
plt.title("Age vs Cholesterol Colored by Heart Attack Risk")
plt.show()
```

## Age vs Cholesterol Colored by Heart Attack Risk



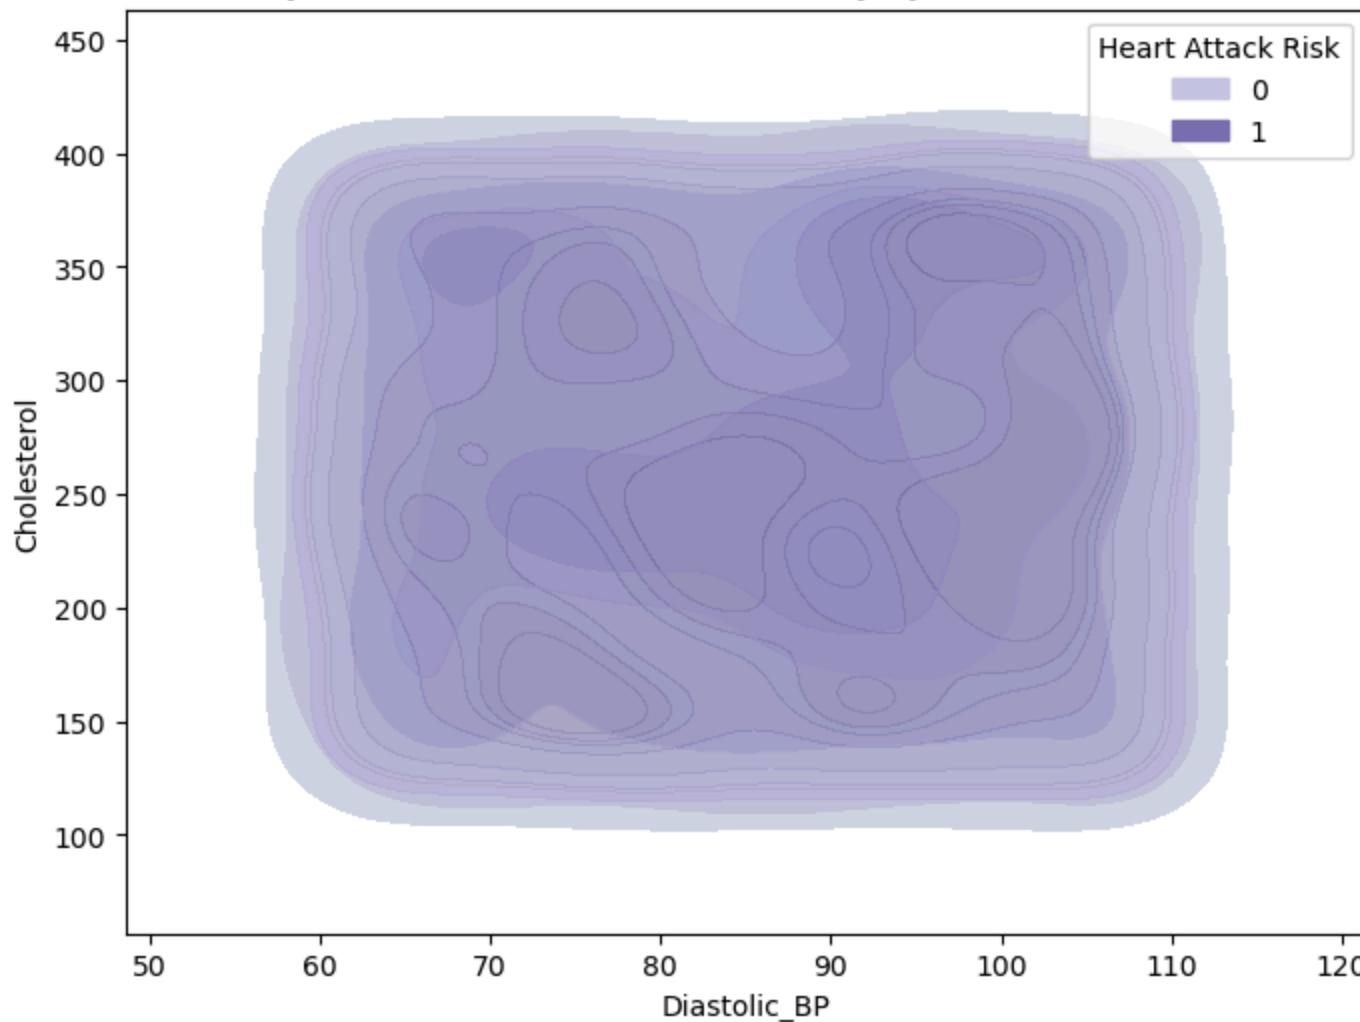
```
In [20]: # Stress Level vs Sleep Hours
plt.figure(figsize=(8,6))
sns.violinplot(data=df, x="Stress Level", y="Sleep Hours Per Day", hue="Heart Attack Risk", split=True, palette="Greens")
plt.title("Stress Level vs Sleep Hours – Distribution by Heart Attack Risk")
plt.show()
```

## Stress Level vs Sleep Hours — Distribution by Heart Attack Risk



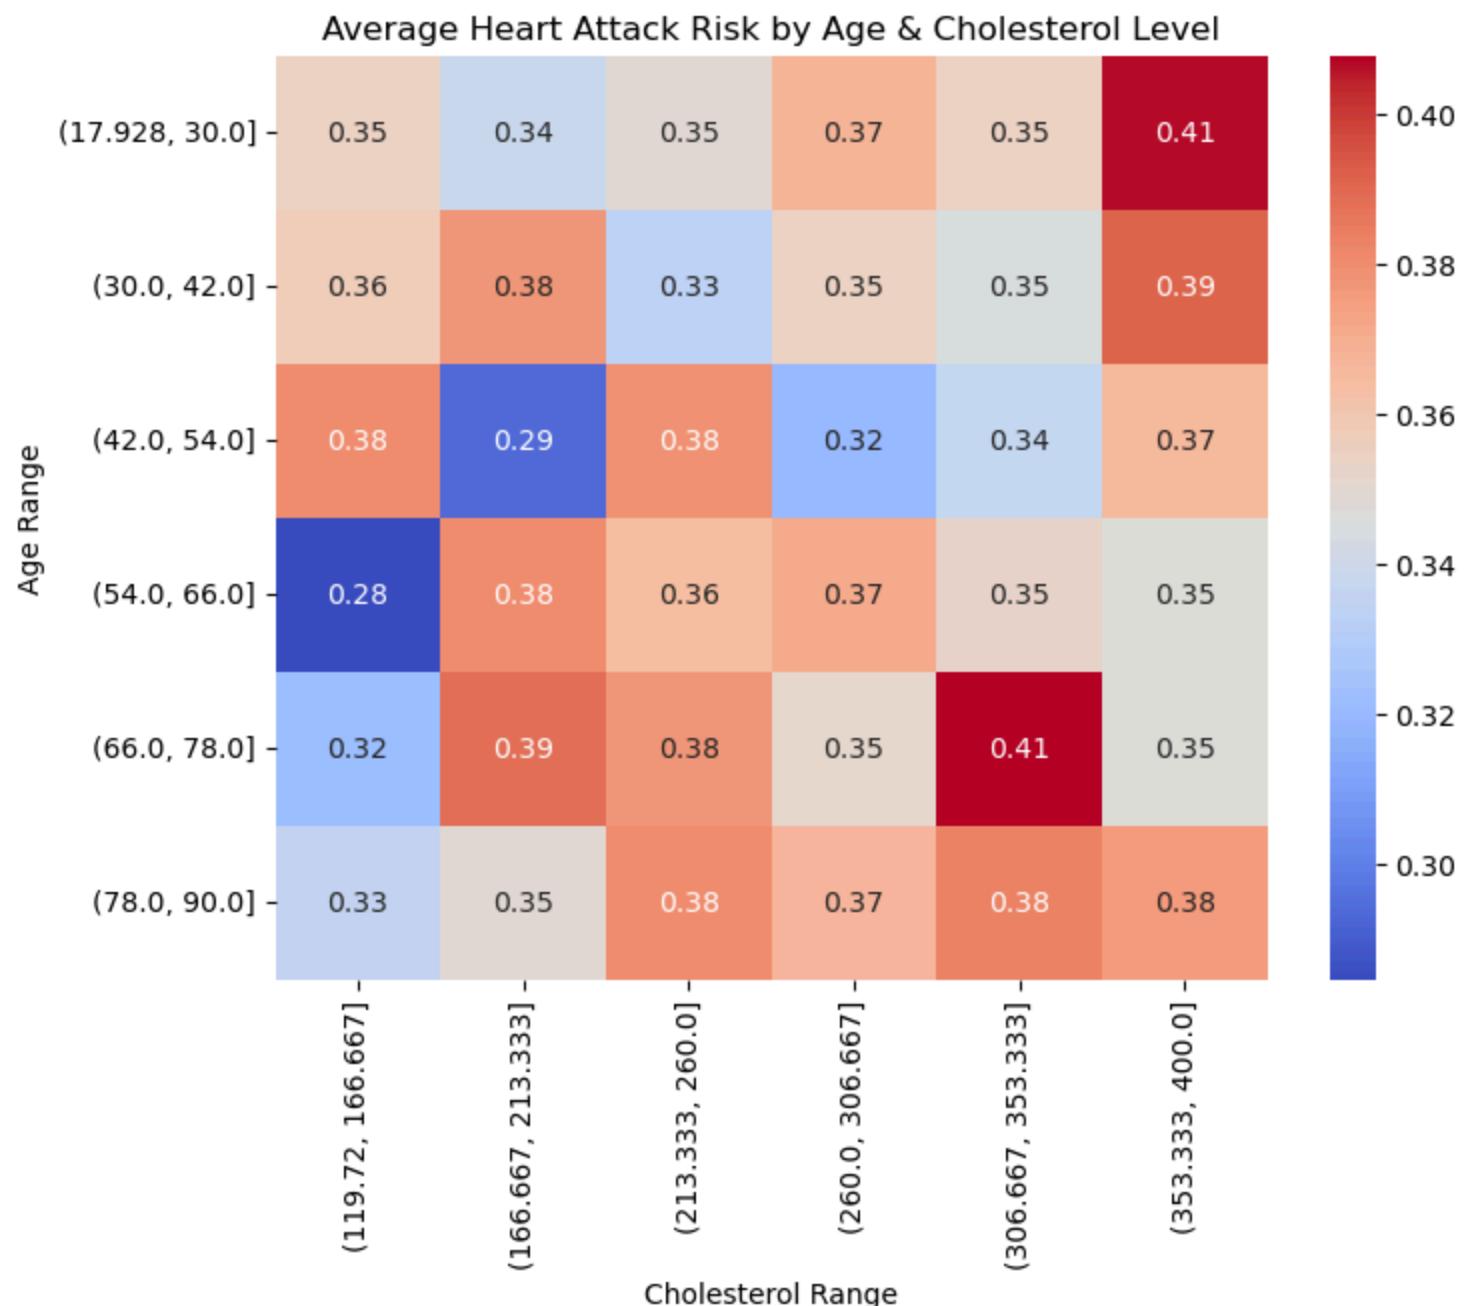
```
In [21]: # Systolic BP vs Cholesterol
plt.figure(figsize=(8,6))
sns.kdeplot(data=df, x="Diastolic_BP", y="Cholesterol", hue="Heart Attack Risk", fill=True, thresh=0.05, alpha=0.6,
)
plt.title("Systolic BP vs Cholesterol – Density by Heart Attack Risk")
plt.show()
```

## Systolic BP vs Cholesterol — Density by Heart Attack Risk



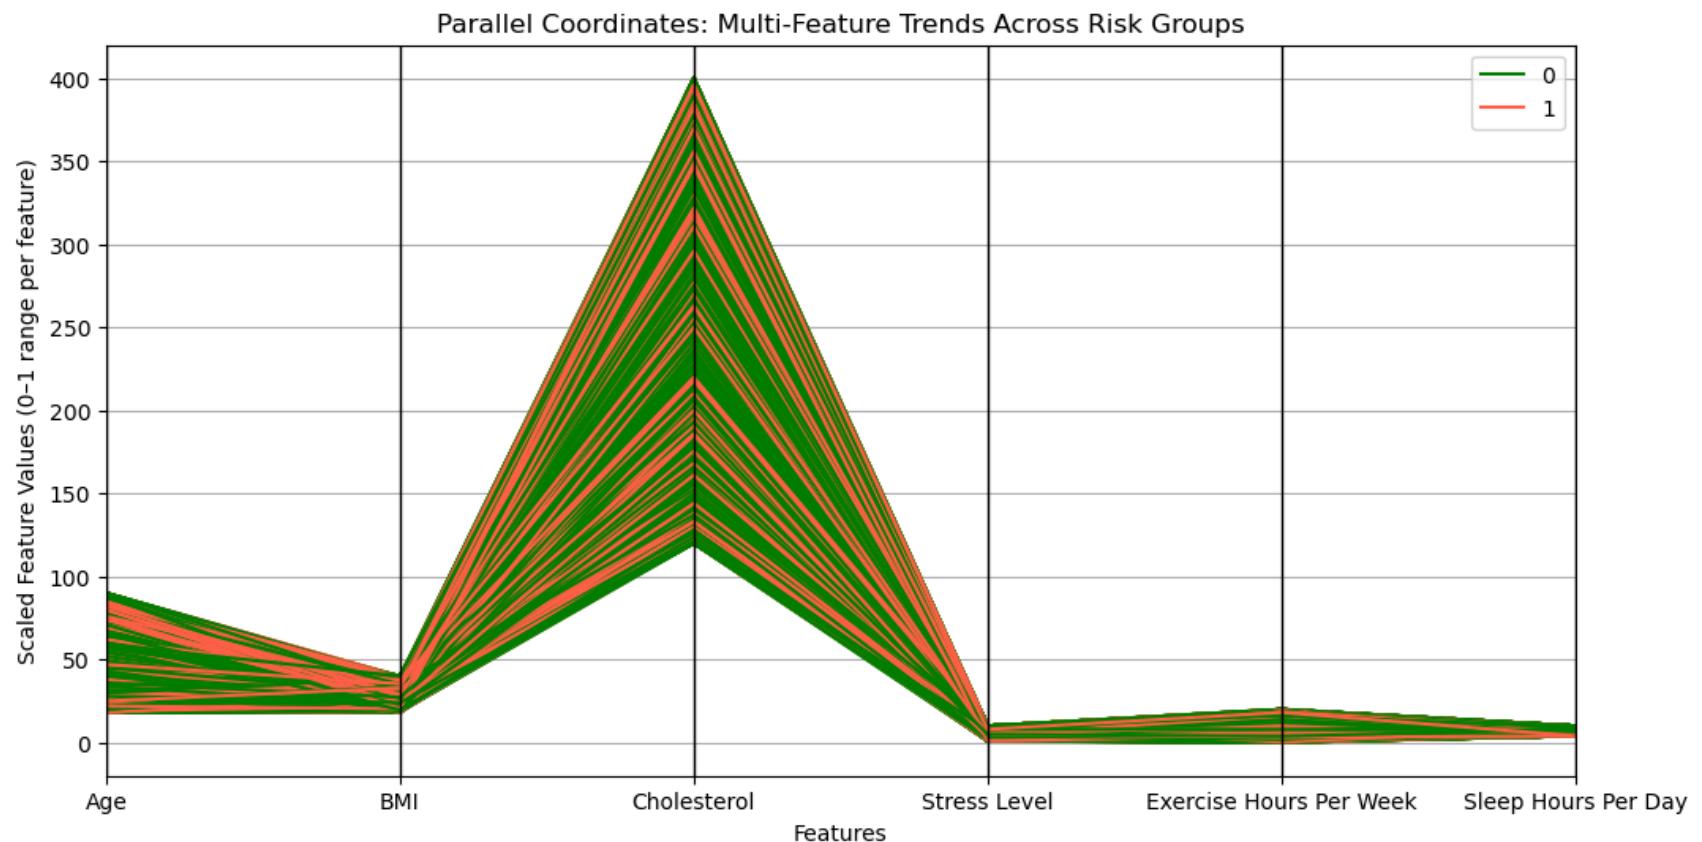
```
In [22]: # Average Heart Attack Risk by Age & Cholesterol Level
plt.figure(figsize=(8,6))
pivot = df.pivot_table(values="Heart Attack Risk",
index=pd.cut(df["Age"], bins=6),
columns=pd.cut(df["Cholesterol"], bins=6),
aggfunc="mean")
sns.heatmap(pivot, cmap="coolwarm", annot=True, fmt=".2f")
plt.title("Average Heart Attack Risk by Age & Cholesterol Level")
plt.xlabel("Cholesterol Range")
```

```
plt.ylabel("Age Range")
plt.show()
```



In [23]:

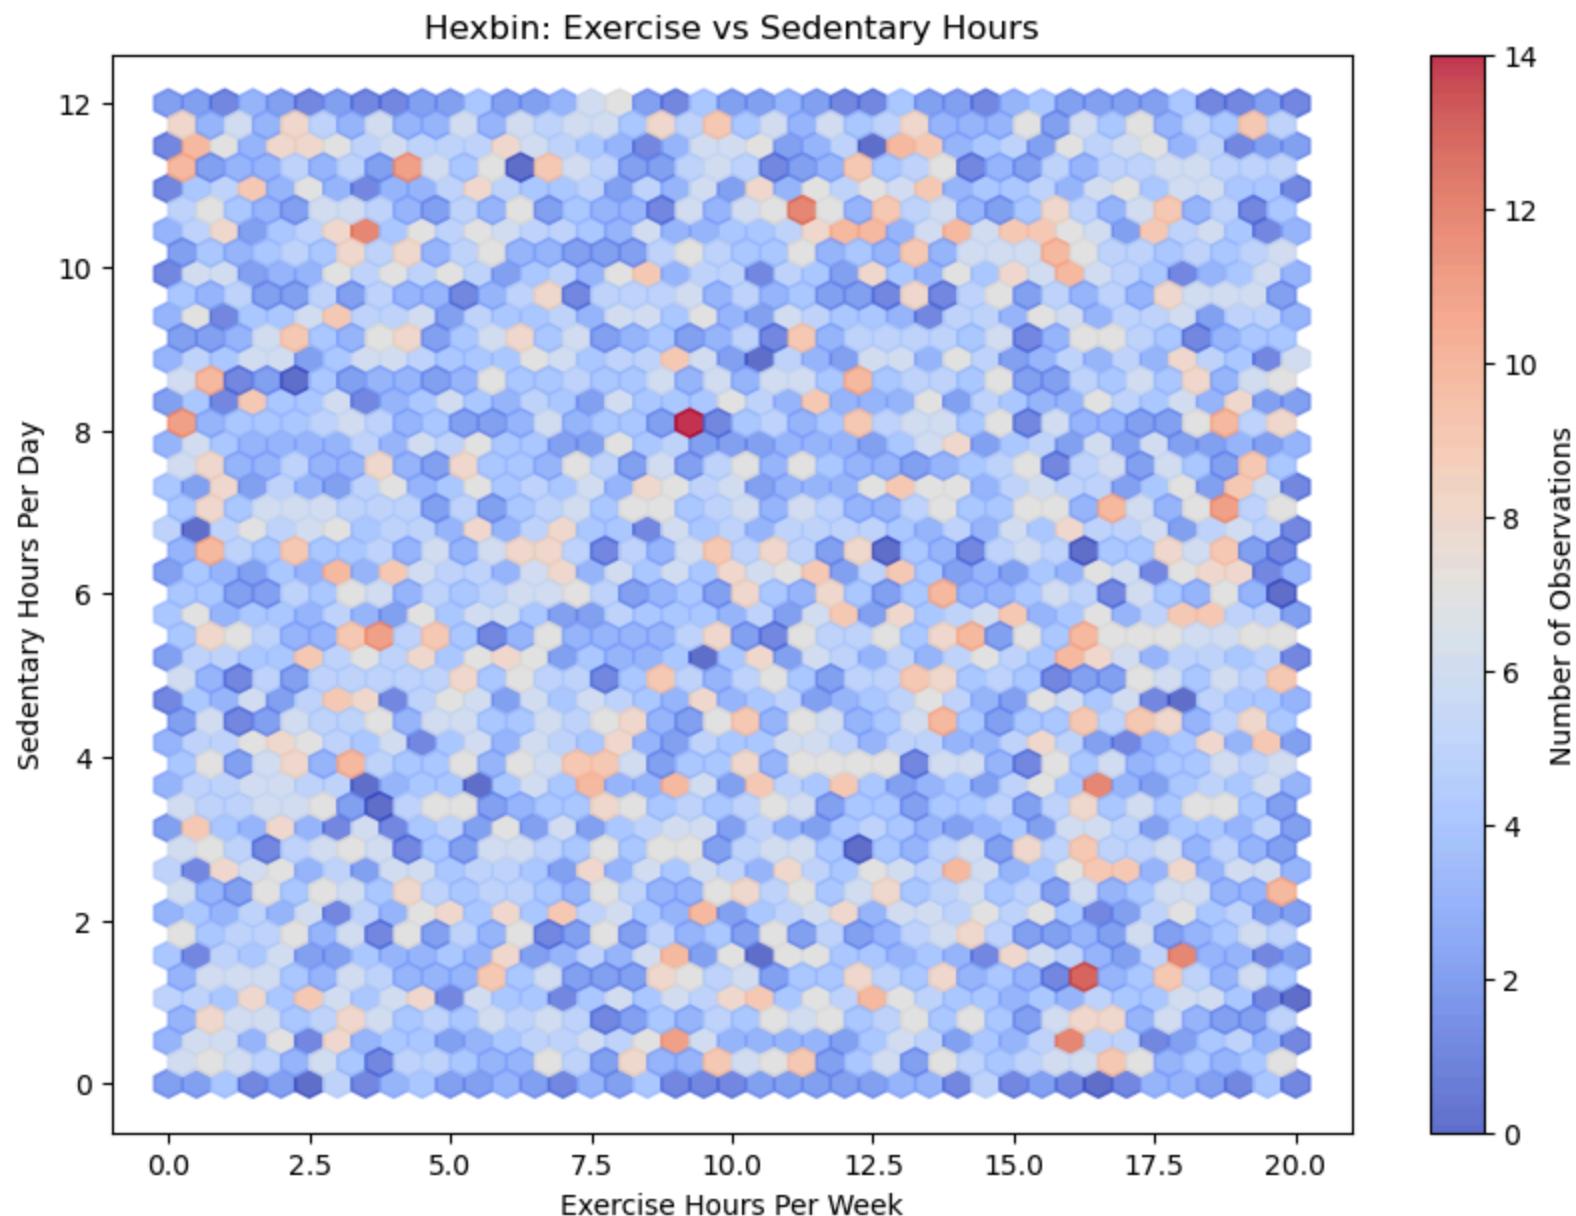
```
# Multi-Feature Trends Across Risk Groups
from pandas.plotting import parallel_coordinates
selected_cols = ["Age", "BMI", "Cholesterol", "Stress Level", "Exercise Hours Per Week", "Sleep Hours Per Day", "Heart Attack Risk"]
plt.figure(figsize=(12,6))
parallel_coordinates(df[selected_cols], "Heart Attack Risk", color=["green", "tomato"])
plt.title("Parallel Coordinates: Multi-Feature Trends Across Risk Groups")
plt.ylabel("Scaled Feature Values (0-1 range per feature)")
plt.xlabel("Features")
plt.show()
```



In [24]:

```
# Hexbin: Exercise vs Sedentary Hours
plt.figure(figsize=(10,7))
plt.hexbin(df["Exercise Hours Per Week"], df["Sedentary Hours Per Day"], gridsize=40, cmap="coolwarm", alpha=0.8)
plt.colorbar(label="Number of Observations")
```

```
plt.title("Hexbin: Exercise vs Sedentary Hours")
plt.xlabel("Exercise Hours Per Week")
plt.ylabel("Sedentary Hours Per Day")
plt.show()
```

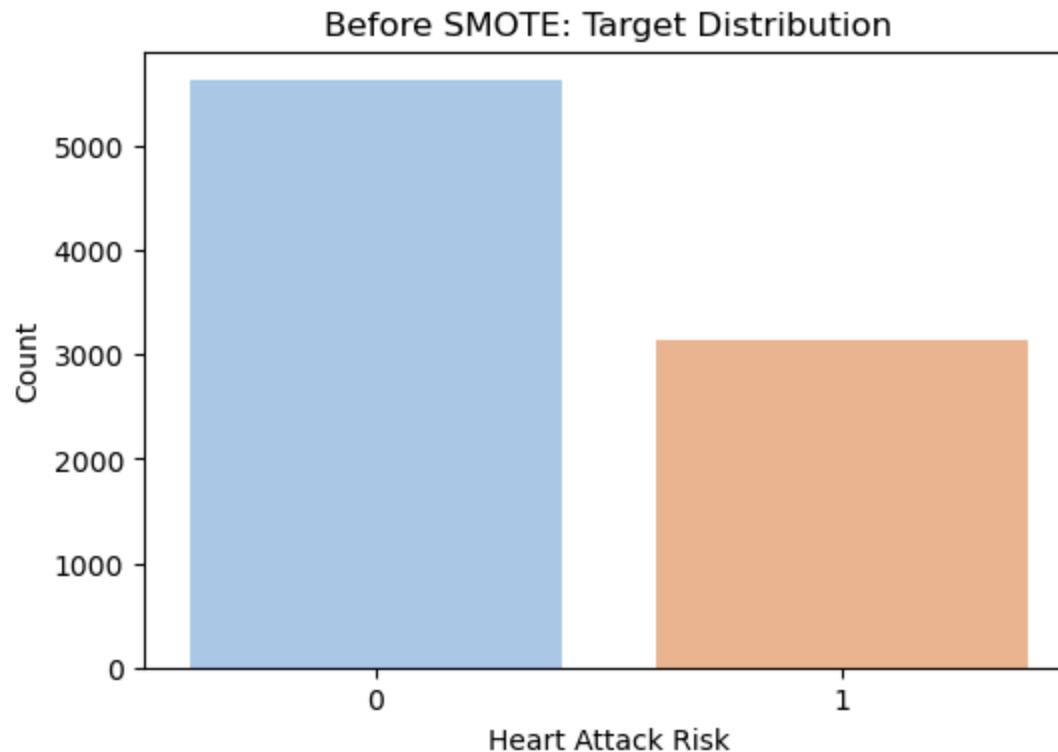


# APPLY SMOTE AND VISUALIZE TARGET BALANCE

## BEFORE SMOTE - CHECK CLASS DISTRIBUTION

```
In [25]: from imblearn.over_sampling import SMOTE
from collections import Counter
X = df.drop(columns=["Heart Attack Risk"])
y = df["Heart Attack Risk"]
X = pd.get_dummies(X, drop_first=True)
print("Original class distribution:")
print(Counter(y))
plt.figure(figsize=(6,4))
sns.countplot(x=y, palette="pastel")
plt.title("Before SMOTE: Target Distribution")
plt.xlabel("Heart Attack Risk")
plt.ylabel("Count")
plt.show()
```

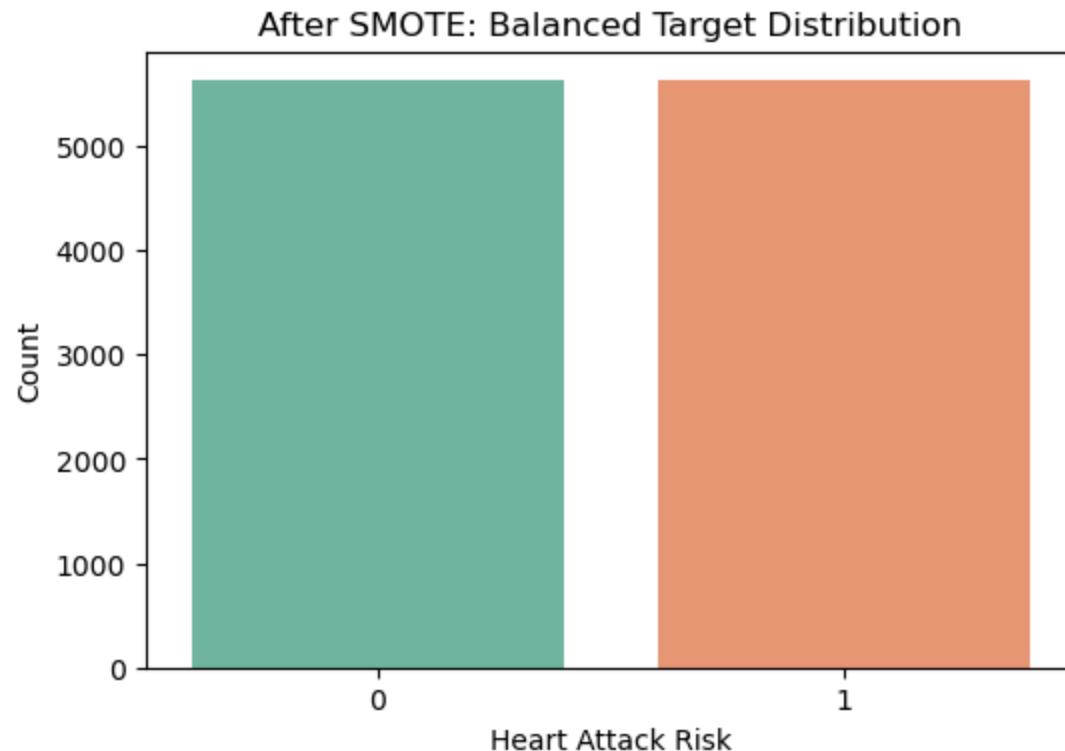
Original class distribution:  
Counter({0: 5624, 1: 3139})



## AFTER SMOTE — CHECK NEW BALANCE

```
In [26]: # Apply SMOTE
smote = SMOTE(random_state=42)
X_res, y_res = smote.fit_resample(X, y)
print("\n After SMOTE class distribution:")
print(Counter(y_res))
plt.figure(figsize=(6,4))
sns.countplot(x=y_res, palette="Set2")
plt.title("After SMOTE: Balanced Target Distribution")
plt.xlabel("Heart Attack Risk")
plt.ylabel("Count")
plt.show()
```

```
After SMOTE class distribution:
Counter({0: 5624, 1: 5624})
```



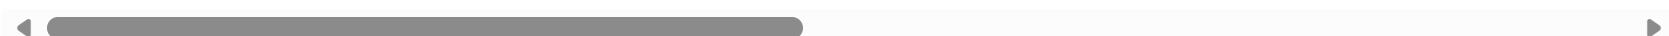
## CREATE A NEW BALANCED DATAFRAME

```
In [27]: df_balanced = pd.concat([pd.DataFrame(X_res), pd.Series(y_res, name="Heart Attack Risk")], axis=1)  
df_balanced
```

Out[27]:

	Age	Cholesterol	Heart Rate	Diabetes	Family History	Smoking	Obesity	Alcohol Consumption	Exercise Hours Per Week	Previous Heart Problems	...	Country_New Zealand	C
0	67	208	72	0	0	1	0	0	4.168189	0	...	False	
1	21	389	98	1	1	1	1	1	1.813242	1	...	False	
2	21	324	72	1	0	0	0	0	2.078353	1	...	False	
3	84	383	73	1	1	1	0	1	9.828130	1	...	False	
4	66	318	93	1	1	1	1	0	5.804299	1	...	False	
...	...	...	...	...	...	...	...	...	...	...	...	...	...
11243	43	201	100	1	0	1	0	0	2.202700	0	...	False	
11244	65	303	74	0	0	1	1	0	10.382115	1	...	False	
11245	25	363	62	1	0	1	1	0	9.892264	0	...	False	
11246	24	293	96	0	0	0	0	0	11.497022	1	...	False	
11247	60	312	58	1	0	1	0	1	9.960904	0	...	False	

11248 rows × 48 columns



## PERFORMING FEATURE SELECTION USING DIFFERENT METHODS FOR SELECTION ON THE NEW DATAFRAME AND TRAINING DIFFERENT MODELS WITH THE SELECTED FEATURES WITH HYPERPARAMETER TUNING

## IMPORTING DEPENDENCIES

In [28]:

```
# Imports
from sklearn.model_selection import train_test_split, GridSearchCV, StratifiedKFold
from sklearn.preprocessing import MinMaxScaler
from sklearn.feature_selection import SelectKBest, chi2, RFE, mutual_info_classif
from sklearn.linear_model import LogisticRegression, Lasso
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from catboost import CatBoostClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

## PREPARING DATA FOR NORMALIZING AND TRAINING

In [29]:

```
# Split data
X = df_balanced.drop("Heart Attack Risk", axis=1)
y = df_balanced["Heart Attack Risk"]

scaler = MinMaxScaler()
X_scaled = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, stratify=y, random_state=42)
```

## MODELS DEFINITIONN AND PARAMETERS FOR TURNING

In [30]:

```
models_params = {
    "LogisticRegression": (LogisticRegression(max_iter=2000), {
        "C": [0.01, 0.1, 1, 10],
        "solver": ["liblinear", "lbfgs"]
    }),
    "DecisionTreeClassifier": (DecisionTreeClassifier(random_state=42), {
        "max_depth": [3, 5, 10, None],
        "min_samples_split": [2, 5, 10],
        "criterion":['gini', 'entropy']
    }),
    "XGBoost": (XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42), {
        "n_estimators": [100, 200],
        "learning_rate": [0.01, 0.1, 0.2],
        "max_depth": [3, 5, 7]
    })
}
```

```
}),
"CatBoost": (CatBoostClassifier(verbose=0, random_state=42), {
    "iterations": [100, 200],
    "learning_rate": [0.01, 0.1, 0.2],
    "depth": [3, 5, 7]
}),
"RandomForest": (RandomForestClassifier(random_state=42), {
    "n_estimators": [100, 200],
    "min_samples_split": [2, 5],
    "max_depth": [5, 10, 15]
})
}
```

## DEFINING FEATURE SELECTION METHODS

```
In [31]: feature_selectors = {}

# 1. Univariate (SelectKBest)
skb = SelectKBest(score_func=chi2, k=10)
skb.fit(X_scaled, y)
feature_selectors['SelectKBest'] = X_scaled.columns[skb.get_support()]

# 2. RFE with Logistic Regression
rfe = RFE(LogisticRegression(max_iter=1000), n_features_to_select=10)
rfe.fit(X_scaled, y)
feature_selectors['RFE'] = X_scaled.columns[rfe.support_]

# 3. Random Forest Importance
rf = RandomForestClassifier(random_state=42)
rf.fit(X_scaled, y)
rf_features = X_scaled.columns[np.argsort(rf.feature_importances_)[:-10:]]
feature_selectors['RandomForest_Importance'] = rf_features

# 4. L1 Regularization (Lasso)
lasso = Lasso(alpha=0.001, max_iter=10000)
lasso.fit(X_scaled, y)
lasso_features = X_scaled.columns[np.abs(lasso.coef_) > 1e-4]
feature_selectors['L1_Lasso'] = lasso_features

# 5. Mutual Information
mi = mutual_info_classif(X_scaled, y)
```

```
mi_features = X_scaled.columns[np.argsort(mi)[-10:]]
feature_selectors['Mutual_Info'] = mi_features
```

## TRAINING AND EVALUATING MODELS

```
In [32]: results = []
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
for method, features in feature_selectors.items():
    print(f"\n Feature Selection: {method}")
    X_train_fs, X_test_fs = X_train[features], X_test[features]

    for name, (model, params) in models_params.items():
        print(f"Training {name} ...")
        grid = GridSearchCV(model, params, cv=cv, scoring='accuracy', n_jobs=-1)
        grid.fit(X_train_fs, y_train)
        best_model = grid.best_estimator_
        y_pred = best_model.predict(X_test_fs)
        acc = accuracy_score(y_test, y_pred)
        pre = precision_score(y_test, y_pred)
        recall = recall_score(y_test, y_pred)
        f1 = f1_score(y_test, y_pred)

        results.append({
            "Feature_Selection": method,
            "Model": name,
            "Best_Params": grid.best_params_,
            "Accuracy": acc,
            "Precision": pre,
            "Recall": recall,
            "F1_Score": f1
        })
```

```
Feature Selection: SelectKBest
Training LogisticRegression ...
Training DecisionTreeClassifier ...
Training XGBoost ...
Training CatBoost ...
Training RandomForest ...
```

```
Feature Selection: RFE
Training LogisticRegression ...
Training DecisionTreeClassifier ...
Training XGBoost ...
Training CatBoost ...
Training RandomForest ...
```

```
Feature Selection: RandomForest_Importance
Training LogisticRegression ...
Training DecisionTreeClassifier ...
Training XGBoost ...
Training CatBoost ...
Training RandomForest ...
```

```
Feature Selection: L1_Lasso
Training LogisticRegression ...
Training DecisionTreeClassifier ...
Training XGBoost ...
Training CatBoost ...
Training RandomForest ...
```

```
Feature Selection: Mutual_Info
Training LogisticRegression ...
Training DecisionTreeClassifier ...
Training XGBoost ...
Training CatBoost ...
Training RandomForest ...
```

## COMPARING RESULTS

```
In [33]: results_df = pd.DataFrame(results)
results_df = results_df.sort_values(by="Accuracy", ascending=False).reset_index(drop=True)
```

```
print("\n Model Performance Summary:")
results_df
```

Model Performance Summary:

Out[33]:

	Feature_Selection	Model	Best_Params	Accuracy	Precision	Recall	F1_Score
0	L1_Lasso	LogisticRegression	{'C': 10, 'solver': 'liblinear'}	0.692889	0.800554	0.513778	0.625880
1	L1_Lasso	RandomForest	{'max_depth': 15, 'min_samples_split': 2, 'n_e...}	0.684889	0.729581	0.587556	0.650911
2	L1_Lasso	XGBoost	{'learning_rate': 0.1, 'max_depth': 3, 'n_esti...}	0.678222	0.714898	0.592889	0.648202
3	L1_Lasso	CatBoost	{'depth': 5, 'iterations': 100, 'learning_rate...}	0.665333	0.706667	0.565333	0.628148
4	Mutual_Info	XGBoost	{'learning_rate': 0.2, 'max_depth': 7, 'n_esti...}	0.644889	0.647378	0.636444	0.641865
5	RandomForest_Importance	CatBoost	{'depth': 7, 'iterations': 200, 'learning_rate...}	0.643556	0.632486	0.685333	0.657850
6	Mutual_Info	CatBoost	{'depth': 5, 'iterations': 200, 'learning_rate...}	0.643111	0.657227	0.598222	0.626338
7	SelectKBest	CatBoost	{'depth': 5, 'iterations': 100, 'learning_rate...}	0.640889	0.679502	0.533333	0.597610
8	SelectKBest	RandomForest	{'max_depth': 5, 'min_samples_split': 2, 'n_es...}	0.634222	0.650398	0.580444	0.613434
9	SelectKBest	XGBoost	{'learning_rate': 0.1, 'max_depth': 5, 'n_esti...}	0.631556	0.671296	0.515556	0.583208
10	SelectKBest	DecisionTreeClassifier	{'criterion': 'entropy', 'max_depth': 10, 'min...}	0.628889	0.676829	0.493333	0.570694
11	RandomForest_Importance	RandomForest	{'max_depth': 15, 'min_samples_split': 2, 'n_e...}	0.628000	0.613386	0.692444	0.650522
12	RandomForest_Importance	XGBoost	{'learning_rate': 0.2, 'max_depth': 7, 'n_esti...}	0.627556	0.623388	0.644444	0.633741
13	Mutual_Info	RandomForest	{'max_depth': 15, 'min_samples_split': 2, 'n_e...}	0.625333	0.628885	0.611556	0.620099
14	SelectKBest	LogisticRegression	{'C': 1, 'solver': 'liblinear'}	0.622667	0.625683	0.610667	0.618084

	Feature_Selection	Model	Best_Params	Accuracy	Precision	Recall	F1_Score
15	L1_Lasso	DecisionTreeClassifier	{'criterion': 'entropy', 'max_depth': 5, 'min_...}	0.592889	0.616760	0.490667	0.546535
16	Mutual_Info	LogisticRegression	{'C': 10, 'solver': 'liblinear'}	0.586222	0.586607	0.584000	0.585301
17	RandomForest_Importance	DecisionTreeClassifier	{'criterion': 'entropy', 'max_depth': None, 'm...}	0.581778	0.576285	0.617778	0.596311
18	Mutual_Info	DecisionTreeClassifier	{'criterion': 'entropy', 'max_depth': 10, 'min...}	0.576444	0.576649	0.575111	0.575879
19	RFE	CatBoost	{'depth': 3, 'iterations': 100, 'learning_rate...}	0.573778	0.564441	0.646222	0.602569
20	RFE	XGBoost	{'learning_rate': 0.1, 'max_depth': 3, 'n_esti...}	0.573778	0.564441	0.646222	0.602569
21	RFE	DecisionTreeClassifier	{'criterion': 'gini', 'max_depth': 10, 'min_sa...}	0.573778	0.564441	0.646222	0.602569
22	RFE	LogisticRegression	{'C': 0.01, 'solver': 'liblinear'}	0.573778	0.564441	0.646222	0.602569
23	RFE	RandomForest	{'max_depth': 5, 'min_samples_split': 2, 'n_es...}	0.568444	0.593674	0.433778	0.501284
24	RandomForest_Importance	LogisticRegression	{'C': 0.01, 'solver': 'liblinear'}	0.513333	0.512755	0.536000	0.524120